


```

1 !pip install --upgrade transformers datasets accelerate
2 from google.colab import drive
3 drive.mount('/content/drive')
4 from sklearn.metrics import accuracy_score, precision_recall_fscore_support
5 import transformers
6 import pandas as pd
7 from datasets import Dataset
8 from transformers import BertTokenizer, BertForSequenceClassification, Trainer,
  TrainingArguments
9 from sklearn.model_selection import train_test_split
10 import torch


```

 Show hidden output

```

1 def compute_metrics(pred):
2     labels = pred.label_ids
3     preds = pred.predictions.argmax(-1)
4     acc = accuracy_score(labels, preds)
5     precision, recall, f1, _ = precision_recall_fscore_support(labels, preds, average='binary')
6     return {
7         'accuracy': acc,
8         'precision': precision,
9         'recall': recall,
10        'f1': f1
11    }

```

 Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True).

```


1 # Only use the input text and hate label
2 df = pd.read_csv('/content/drive/MyDrive/train_with_topics.csv')
3 df = df[['text', 'topic_label', 'hate_label']].dropna()
4 df['hate_label'] = df['hate_label'].astype(int)
5
6 # Split into train/test
7 from sklearn.model_selection import train_test_split
8 train_df, val_df = train_test_split(df, test_size=0.1, random_state=42)
9
10 # Convert to Hugging Face dataset
11 train_dataset = Dataset.from_pandas(train_df)
12 val_dataset = Dataset.from_pandas(val_df)
13 full_dataset = Dataset.from_pandas(df)

```

```

1 tokenizer = BertTokenizer.from_pretrained("bert-base-uncased")
2
3 def tokenize(example):
4     return tokenizer(example['text'], truncation=True, padding="max_length", max_length=128)
5
6 train_dataset = train_dataset.map(tokenize, batched=True)
7 val_dataset = val_dataset.map(tokenize, batched=True)
8 full_dataset = full_dataset.map(tokenize, batched=True)
9
10 train_dataset = train_dataset.rename_column("hate_label", "labels")
11 val_dataset = val_dataset.rename_column("hate_label", "labels")
12 full_dataset = full_dataset.rename_column("hate_label", "labels")
13
14 train_dataset.set_format("torch")
15 val_dataset.set_format("torch")
16 full_dataset.set_format("torch")
17

```

 Map: 100% 13927/13927 [00:08<00:00, 1621.94 examples/s]

Map: 100% 1548/1548 [00:00<00:00, 1679.21 examples/s]

Map: 100% 15475/15475 [00:09<00:00, 1727.65 examples/s]

```

1 # Load model
2 model = BertForSequenceClassification.from_pretrained("bert-base-uncased", num_labels=2)
3
4 # Define training arguments
5 training_args = TrainingArguments(
6     output_dir="/content/drive/MyDrive/hate_model_basic_full_dataset",
7     save_strategy="epoch",
8     per_device_train_batch_size=8,
9     per_device_eval_batch_size=8,
10    num_train_epochs=10,
11    logging_dir="/content/logs"
12 )
13
14 trainer = Trainer(
15     model=model,
16     args=training_args,
17     train_dataset=train_dataset,
18     eval_dataset=val_dataset,
19     tokenizer=tokenizer,
20     compute_metrics=compute_metrics
21 )
22
23 # Train and evaluate
24 trainer.train()
25 trainer.evaluate()
26

```

Some weights of BertForSequenceClassification were not initialized at bert-base-uncased and are newly initialized: ['classifier.bias', 'classifier.weight']. You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference.

```
<ipython-input-20-90d7fb882616>:14: FutureWarning: `tokenizer` is deprecated and will be removed in version 5.0.0 for `Trainer.__init__`. Use `processing_class` instead:
trainer = Trainer(
[13849/17410 12:47 < 03:17, 18.05 it/s, Epoch 7.95/10]
```

Step	Training Loss
------	---------------

500	0.461300
1000	0.450400
1500	0.423300
2000	0.389100
2500	0.354800
3000	0.357800
3500	0.342500
4000	0.248700
4500	0.260900
5000	0.264600
5500	0.209900
6000	0.188000
6500	0.159400
7000	0.193800
7500	0.100000
8000	0.117100
8500	0.120300
9000	0.095100
9500	0.068600
10000	0.060700
10500	0.063500
11000	0.044400
11500	0.058300
12000	0.055000
12500	0.040100
13000	0.033100
13500	0.032600

[17410/17410 16:06, Epoch 10/10]

Step	Training Loss
------	---------------

500	0.461300
1000	0.450400
1500	0.423300
2000	0.389100
2500	0.354800
3000	0.357800
3500	0.342500
4000	0.248700
4500	0.260900
5000	0.264600
5500	0.209900
6000	0.188000
6500	0.159400
7000	0.193800
7500	0.100000
8000	0.117100
8500	0.120300
9000	0.095100
9500	0.068600
10000	0.060700
10500	0.063500
11000	0.044400
11500	0.058300
12000	0.055000
12500	0.040100
13000	0.033100
13500	0.032600
14000	0.040500
14500	0.020700
15000	0.024100
15500	0.026300
16000	0.012700
16500	0.016100

```
17000 0.020200 [194/194 00:03]
{'eval_loss': 1.321150779724121,
 'eval_accuracy': 0.8275193798449613,
 'eval_precision': 0.625,
 'eval_recall': 0.5539358600583091,
 'eval_f1': 0.5873261205564142,
 'eval_runtime': 3.1807,
 'eval_samples_per_second': 486.685,
 'eval_steps_per_second': 60.993,
 'epoch': 10.0}
```

```
1 # Running model in-domain on Twitter dataset
2 # Load tokenizer
3 tokenizer = BertTokenizer.from_pretrained("bert-base-uncased")
4
5 # Load test CSV
6 df = pd.read_csv("/content/drive/MyDrive/train.csv")
7
8 # Batch size for prediction
9 batch_size = 32
10
11 # Get the device
12 device = "cuda" if torch.cuda.is_available() else "cpu"
13 model.to(device) # Ensure model is on the correct device
14
15 # Store predictions
16 all_preds = []
17 all_probs = []
18
19 # Iterate over the data in batches
20 for i in range(0, len(df), batch_size):
21     batch = df[i : i + batch_size]
22
23     test_encodings = tokenizer(
24         batch["text"].tolist(),
25         truncation=True,
26         padding=True,
27         max_length=128,
28         return_tensors="pt",
29     ).to(device)
30
31     model.eval()
32     with torch.no_grad():
33         outputs = model(**test_encodings)
34         probs = torch.softmax(outputs.logits, dim=1)
35         preds = torch.argmax(probs, dim=1)
36
37     all_preds.extend(preds.cpu().numpy())
38     all_probs.extend(probs.cpu().numpy().tolist())
39
40 # Add predictions to DataFrame
41 df["predicted_label"] = all_preds
42 df["predicted_prob"] = all_probs
43
44 # Calculate accuracy
45 correct = (df["predicted_label"] == df["hate_label"]).sum()
46 accuracy = correct / len(df)
47
48 # Print results
49 print(df[["text", "hate_label", "predicted_label"]].head())
50 print(f"Accuracy: {accuracy:.4f}")
51
```

```

text hate_label \
0 The trans women reading this tweet right now i... 0
1 9) uhhhh i like being lgbt a lot. i feel proud... 0
2 @terryelaine1 @UKLabour Why do 3.8 million #5... 0
3 I said it yesterday, I knew this is about to g... 0
4 White Small Little Invisible Clits Are A Disgr... 1

predicted_label
0 0
1 0
2 0
3 0
4 1
Accuracy: 0.9899
```

```
1 # Running model out of domain on Reddit dataset
2
3 # Load tokenizer
4 tokenizer = BertTokenizer.from_pretrained("bert-base-uncased")
5
6 # Load test CSV
7 df = pd.read_csv("/content/drive/MyDrive/test_reddit.csv")
8
9 # Get the device
10 device = "cuda" if torch.cuda.is_available() else "cpu"
11 model.to(device) # Ensure model is on the correct device
12
13 # Batch size for prediction
14 batch_size = 32
15
16 # Store predictions
17 all_preds = []
18 all_probs = []
19
20 # Iterate over the data in batches
21 for i in range(0, len(df), batch_size):
22     batch = df[i : i + batch_size]
23
24     test_encodings = tokenizer(
25         batch["text"].tolist(),
26         truncation=True,
27         padding=True,
28         max_length=128,
29         return_tensors="pt",
30     ).to(device)
```

```

31
32     model.eval()
33     with torch.no_grad():
34         outputs = model(**test_encodings)
35         probs = torch.softmax(outputs.logits, dim=1)
36         preds = torch.argmax(probs, dim=1)
37
38     all_preds.extend(preds.cpu().numpy())
39     all_probs.extend(probs.cpu().numpy().tolist())
40
41 # Add predictions to DataFrame
42 df["predicted_label"] = all_preds
43 df["predicted_prob"] = all_probs
44
45 # Calculate accuracy
46 correct = (df["predicted_label"] == df["hate_label"]).sum()
47 accuracy = correct / len(df)
48
49 # Print results
50 print(df[["text", "hate_label", "predicted_label"]].head())
51 print(f"Accuracy: {accuracy:.4f}")
52

```

```

35
0  For starters bend over the one in pink and kic...    1
1  Sounds like the kinda wholesome life I'd die ...    0
2  Who the fuck is this insignificant simple mind...    0
3      Fuck off you insufferable retarded faggot.      1
4  Worthless whore, these tits with look nice wit...    0

predicted_label
0          0
1          0
2          1
3          1
4          1
Accuracy: 0.7957

```

```

1 # Running model out of domain on Youtube dataset
2
3 # Load tokenizer
4 tokenizer = BertTokenizer.from_pretrained("bert-base-uncased")
5
6 # Load test CSV
7 df = pd.read_csv("/content/drive/MyDrive/test_youtube.csv")
8
9 # Batch size for prediction
10 batch_size = 32
11
12 # Store predictions
13 all_preds = []
14 all_probs = []
15
16 # Iterate over the data in batches
17 for i in range(0, len(df), batch_size):
18     batch = df[i : i + batch_size]
19
20     test_encodings = tokenizer(
21         batch["text"].tolist(),
22         truncation=True,
23         padding=True,
24         max_length=128,
25         return_tensors="pt",
26     ).to(device)
27
28     model.eval()
29     with torch.no_grad():
30         outputs = model(**test_encodings)
31         probs = torch.softmax(outputs.logits, dim=1)
32         preds = torch.argmax(probs, dim=1)
33
34     all_preds.extend(preds.cpu().numpy())
35     all_probs.extend(probs.cpu().numpy().tolist())
36
37 # Add predictions to DataFrame
38 df["predicted_label"] = all_preds
39 df["predicted_prob"] = all_probs
40
41 # Calculate accuracy
42 correct = (df["predicted_label"] == df["hate_label"]).sum()
43 accuracy = correct / len(df)
44
45 # Print results
46 print(df[["text", "hate_label", "predicted_label"]].head())
47 print(f"Accuracy: {accuracy:.4f}")
48

```

```

35
0  Yes indeed. She sort of reminds me of the elde...    0
1  Question: These 4 broads who criticize America...    0
2  It is about time for all illegals to go back t...    0
3  OMG! The EGO's of these young, young, inexperi...    0
4  Joshua Lelo so you have seen all actors from e...    0

predicted_label
0          0
1          0
2          0
3          0
4          0
Accuracy: 0.7754

```

